

Course: MMDS

Lab 01

Set up Hadoop

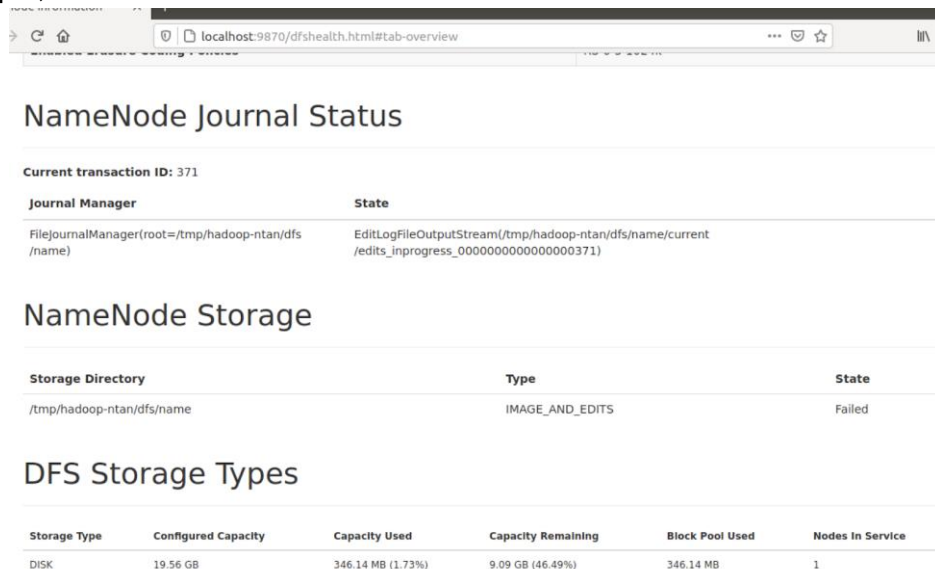
Fill answers of the questions below in the given tables.

Your screenshots must contain commands for required operations.

Question 1:

Set up Hadoop following this [tutorial](#) and then take a screenshot of the hdfs-site ([localhost:9870](#)) which shows information for NameNode Journal Status, NameNode Storage, DFS Storage Types (Overview page).

For example,



The screenshot shows the Hadoop NameNode Overview page. The browser address bar indicates the URL is localhost:9870/dfshealth.html#tab-overview. The page content is as follows:

NameNode Journal Status

Current transaction ID: 371

Journal Manager	State
FileJournalManager(root=/tmp/hadoop-ntan/dfs/name)	EditLogOutputStream(/tmp/hadoop-ntan/dfs/name/current/edits_inprogress_00000000000000000371)

NameNode Storage

Storage Directory	Type	State
/tmp/hadoop-ntan/dfs/name	IMAGE_AND_EDITS	Failed

DFS Storage Types

Storage Type	Configured Capacity	Capacity Used	Capacity Remaining	Block Pool Used	Nodes In Service
DISK	19.56 GB	346.14 MB (1.73%)	9.09 GB (46.49%)	346.14 MB	1

Jan 25 16:35

Course: HK2 | tutorials - Go | Lab 01 - Had | Tutorial 01 - | lectures - Go | references (v | 2324-HK2-M | Namenode in | X

localhost:9870/dfshealth.html#tab-overview

←

→

↻

🔖

🔍

📄

📌

☰

NameNode Journal Status

Current transaction ID: 3

Journal Manager	State
FileJournalManager(root=/tmp/hadoop-trvy/dfs/name)	EditLogOutputStream(/tmp/hadoop-trvy/dfs/name/current/edits_inprogress_0000000000000000003)

NameNode Storage

Storage Directory	Type	State
/tmp/hadoop-trvy/dfs/name	IMAGE_AND_EDITS	Active

DFS Storage Types

Storage Type	Configured Capacity	Capacity Used	Capacity Remaining	Block Pool Used	Nodes In Service
DISK	49.41 GB	24 KB (0%)	36.14 GB (73.15%)	24 KB	1

Jan 25 16:36

Course: HK2 | tutorials - Go | Lab 01 - Had | Tutorial 01 - | lectures - Go | references (v | 2324-HK2-M | Namenode in | X

localhost:9870/dfshealth.html#tab-overview

←

→

↻

🔖

🔍

📄

📌

☰

Hadoop | Overview | Datanodes | Datanode Volume Failures | Snapshot | Startup Progress | Utilities ▾

Overview 'localhost:9000' (active)

Started:	Thu Jan 25 16:31:09 +0700 2024
Version:	3.2.1, rb3cbbb467e22ea829b3808f4b7b01d07e0bf3842
Compiled:	Tue Sep 10 22:56:00 +0700 2019 by rohithsharmaks from branch-3.2.1
Cluster ID:	CID-9340b173-590f-48b9-b2fe-09813914d403
Block Pool ID:	BP-692445599-127.0.1.1-1706175008195

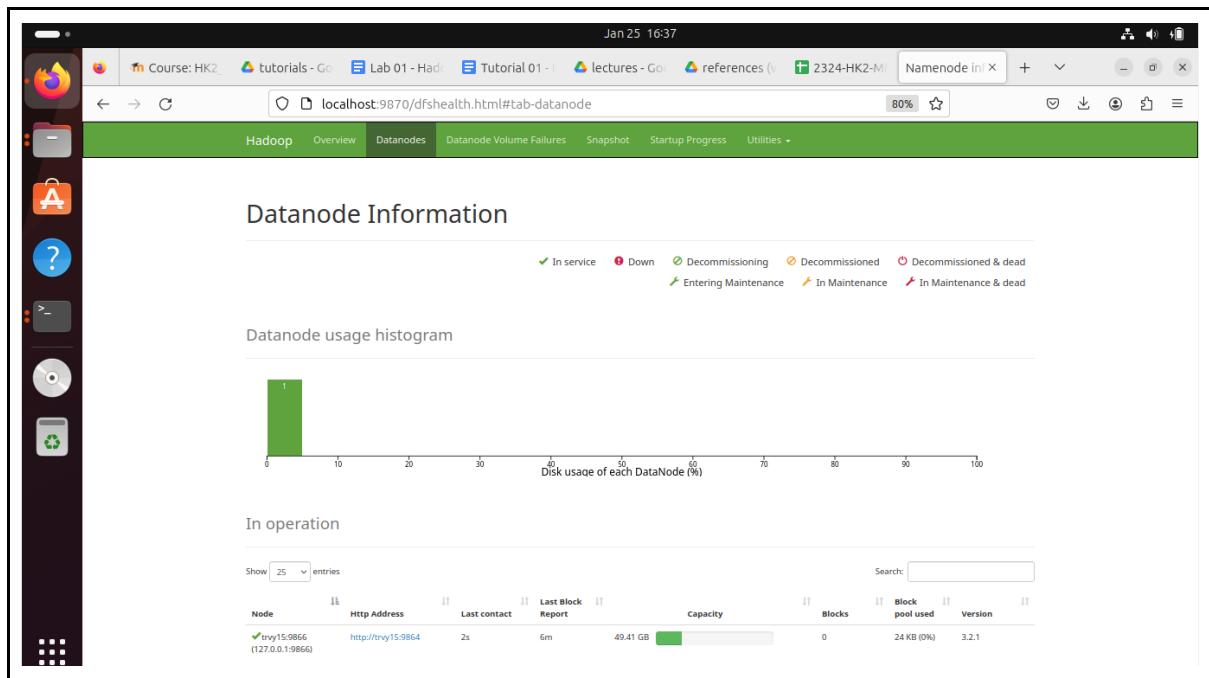
Summary

Security is off.
Safemode is off.

1 files and directories, 0 blocks (0 replicated blocks, 0 erasure coded block groups) = 1 total filesystem object(s).

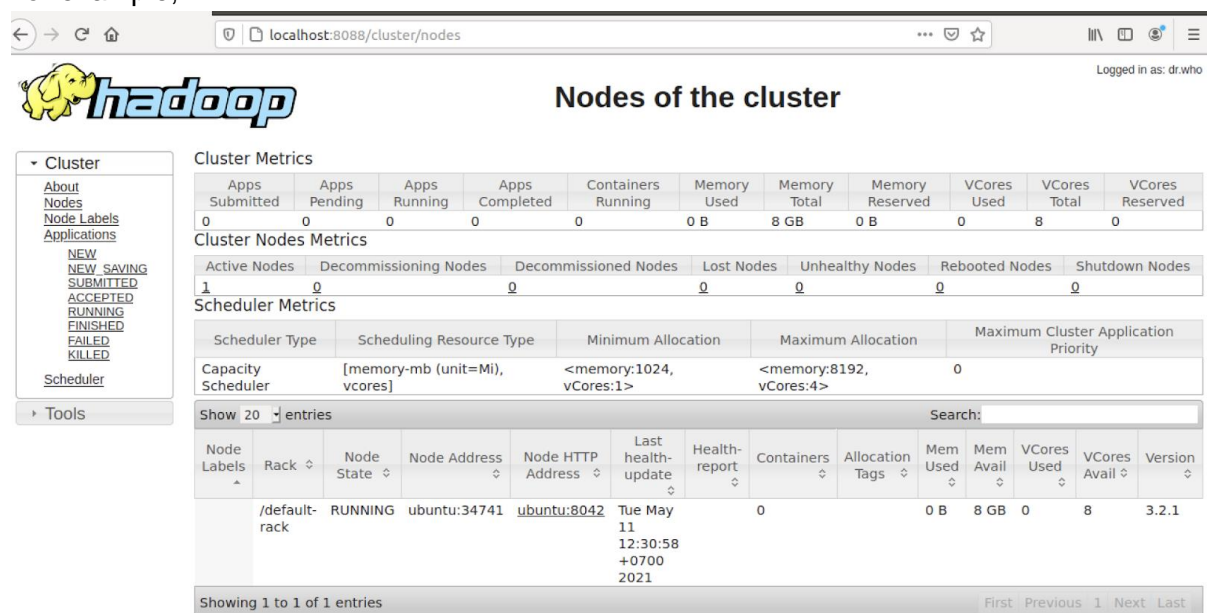
Heap Memory used 110.68 MB of 301.5 MB Heap Memory. Max Heap Memory is 2.39 GB.

Non Heap Memory used 47.33 MB of 49.02 MB Committed Non Heap Memory. Max Non Heap Memory is <unbounded>.



Question 2:


Set up Hadoop following this [tutorial](#) and then take a screenshot of the yarn-site (localhost:8088) which shows information for nodes (<http://localhost:8088/cluster/nodes>). For example,



Jan 25 21:07

Course: HK2_202 × tutorials - Google × Lab 01 - Hadoop × Tutorial 01 - Hadoop × Browsing HDFS × nodes Applications ×

localhost:8088/cluster/apps/nodes 90%



nodes Applications

Logged in as: dr.who

Cluster

About

Nodes

Node Labels

Applications

NEW

NEW SAVING

SUBMITTED

ACCEPTED

RUNNING

FINISHED

FAILED

KILLED

Scheduler

Tools

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved
0	0	0	0	0	0 B	8 GB	0 B	0	8	0

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes	Shutdown Nodes
1	0	0	0	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation	Maximum Cluster Application Priority
Capacity Scheduler	[memory-mb (unit=Mi), vcores]	<memory:1024, vCores:1>	<memory:8192, vCores:4>	0

Failed to read the applications.

Submission Notice

- Export your answer file as pdf
- Rename the pdf following the format:
lab01_<student number>_<Full Name>.pdf
E.g: **lab01_123456_NguyenThanhAn.pdf**
If you have not been assigned a student number yet, then use 123456 instead.
- Careless mistakes in filename, format, question order, etc. are not accepted (0 pts).